



# KMapSolver4D

24-Bit Boolean Minimization

Multicore Performance Report

Variables: 24

CPU Cores: 4

Total Chunks: 256

Date: 2025-12-02

© Stan's Technologies 2025

## EXECUTIVE SUMMARY

---

### EXPERIMENT OVERVIEW

---

This report documents a 24-bit Boolean function minimization using KMapSolver4D with parallel processing across 4 CPU cores.

### CONFIGURATION

---

Variables:	24
K-map variables:	16 (per chunk)
Chunk selector vars:	8
Total chunks:	256
CPU cores:	4
Chunks per core:	64
Random seed:	42

### PERFORMANCE RESULTS

---

Total experiment time:	2.32 hours (139.4 minutes)
Wall-clock time:	2.32 hours
Total CPU time:	8.68 hours
Parallel efficiency:	3.74x speedup
Minimized chunks:	256 / 256
Total terms generated:	4,325,398

### FINAL EXPRESSION ANALYSIS

---

Analyzed chunks:	4,325,398
Unique terms:	0
Sample terms shown:	0
Total literals (sample):	0
Avg lits/term:	0.0
Max term size:	0
Min term size:	0

### SYSTEM INFORMATION

---

Platform:	Windows-11-10.0.26200-SP0
Processor:	Intel64 Family 6 Model 142 Stepping 12, GenuineIntel
Python version:	3.12.10
Execution date:	2025-12-02 18:57:35

### KEY FINDINGS

---

- Successfully minimized 256 chunks in parallel
- Achieved 3.74x speedup with 4 cores
- Generated 4,325,398 product terms total
- Final expression contains 0 unique terms
- Chunked algorithm enables minimization beyond 16-variable limit

### DELIVERABLES

---

- ✓ Performance statistics and timing analysis
- ✓ Final minimized Boolean expression (SOP form)
- ✓ Verilog HDL module with testbench
- ✓ Comprehensive scientific documentation

FINAL MINIMIZED BOOLEAN EXPRESSION (COMPLETE)

---

EXPRESSION FORM: Sum of Products (SOP)

STATISTICS

---

Total unique terms: 0  
Total literals: 0  
Average lits per term: 0.00  
Max term size: 0  
Min term size: 0

EXPRESSION (ALL 0 TERMS)

---

NOTES

---

- This is the COMPLETE expression with all 0 unique terms
- Expression represents bitwise union of all 256 chunks
- Each chunk covers a 16-variable K-map (65,536 entries)
- The 8 extra variables select which chunk to evaluate
- Expression is in minimized SOP (Sum of Products) form
- Total experiment time: 2.32 hours

VERILOG HDL MODULE

---

MODULE: boolean\_function\_24bit  
DESCRIPTION: Hardware description for 24-bit Boolean function

```
// Verilog Module: boolean_function_24bit
// Generated: 2025-12-02 18:57:34
// Variables: 24
// Description: 24-bit Boolean function minimized with KMapSolver4D

module boolean_function_24bit (
    input [23:0] inputs,
    output result
);

// Variable mapping
wire x0 = inputs[0];
wire x1 = inputs[1];
wire x2 = inputs[2];
wire x3 = inputs[3];
wire x4 = inputs[4];
wire x5 = inputs[5];
wire x6 = inputs[6];
wire x7 = inputs[7];
wire x8 = inputs[8];
wire x9 = inputs[9];
wire x10 = inputs[10];
wire x11 = inputs[11];
wire x12 = inputs[12];
wire x13 = inputs[13];
wire x14 = inputs[14];
wire x15 = inputs[15];
wire x16 = inputs[16];
wire x17 = inputs[17];
wire x18 = inputs[18];
wire x19 = inputs[19];
wire x20 = inputs[20];
wire x21 = inputs[21];
wire x22 = inputs[22];
wire x23 = inputs[23];

// Minimized Boolean expression (SOP form)
// COMPLETE expression with all terms
assign result = 1'b0; // Empty expression

endmodule

// Testbench for boolean_function_24bit
module boolean_function_24bit_tb;

reg [23:0] inputs;
wire result;

boolean_function_24bit uut (
    .inputs(inputs),
    .result(result)
);

initial begin
    $dumpfile("boolean_function_24bit.vcd");
    $dumpvars(0, boolean_function_24bit_tb);

    // Test cases
    inputs = 24'b0;
    #10;

    // Add more test vectors here
    inputs = 24'h1;
    #10;

```


VERILOG HDL MODULE (continued)

---


```
inputs = 24'hFFFF;  
#10;  
$finish;  
end  
endmodule
```

# 24-Bit Boolean Minimization: Performance Analysis


**A) Processing Time by Core**



**B) Parallel Efficiency**



**C) Chunk Distribution**



Metric	Value
Total chunks	256
Minimized chunks	256
Terms generated	4,325,398
Unique terms	0
CPU cores	4
Wall-clock time	2.32 hrs
CPU time	8.68 hrs
Speedup	3.74x

## METHODOLOGY & SCIENTIFIC CONCLUSIONS

---

### CHUNKED MINIMIZATION ALGORITHM

---

The KMapSolver4D algorithm divides large Boolean functions into manageable chunks:

#### 1. CHUNKING STRATEGY

- Input: 24-bit Boolean function ( $2^{24} = 16,777,216$  entries)
- Partition: 256 chunks  $\times$  65,536 entries per chunk
- Method: Use 8 variables as chunk selectors, 16 for K-maps
- Processing: Each chunk minimized independently with KMapSolver3D

#### 2. PARALLEL PROCESSING

- Distribution: 256 chunks across 4 cores
- Load balancing: 64 chunks per core
- Synchronization: Independent workers with CSV output
- Aggregation: Bitwise union of chunk expressions

#### 3. MEMORY MANAGEMENT

- No full truth table storage required
- Stream processing: Generate  $\rightarrow$  Minimize  $\rightarrow$  Write  $\rightarrow$  Discard
- Peak memory: ~2-3 GB per core (vs. ~100+ GB for naive approach)

## PERFORMANCE ANALYSIS

---

Wall-clock time: 2.32 hours  
Sequential estimate: 8.68 hours  
Speedup achieved: 3.74x  
Efficiency: 93.4%

The parallel efficiency of 3.74x with 4 cores indicates excellent scalability with minimal overhead.

## EXPRESSION COMPLEXITY

---

Total unique terms: 0  
Sample analyzed: 0 terms  
Literals (sample): 0  
Avg complexity: 0.00 literals/term

The final expression demonstrates the complexity inherent in 24-bit Boolean functions, with 0 unique product terms after minimization.

## KEY FINDINGS

---

#### 1. SCALABILITY

- Chunked algorithm successfully extends K-map minimization to 24 bits
- Memory requirements remain tractable (2-3 GB vs. 100+ GB)
- Processing time scales linearly with chunk count

#### 2. PARALLEL EFFICIENCY

- 4-core parallelization achieved 3.74x speedup
- Independent chunk processing enables embarrassingly parallel workload
- Minimal synchronization overhead

#### 3. SOLUTION QUALITY

- Each chunk optimally minimized with KMapSolver3D
- Final expression is union of all chunk solutions
- 256 / 256 chunks successfully processed

#### 4. PRACTICAL VIABILITY

- 24-bit minimization feasible on commodity hardware
- 2.32 hour runtime acceptable for batch processing
- Verilog output enables direct hardware synthesis

## LIMITATIONS & THREATS TO VALIDITY

---

### INTERNAL VALIDITY

- Random function generation may not reflect real circuits
- Chunk independence assumes no cross-chunk optimization opportunities
- Python runtime overhead included in measurements

### EXTERNAL VALIDITY

- Results specific to this implementation and hardware
- Performance varies with function complexity and distribution
- No comparison with other large-scale minimizers

### CONSTRUCT VALIDITY

- Union of chunk solutions may not be globally optimal
- Literal count as complexity metric has limitations
- Speedup calculation assumes perfect sequential baseline

## RECOMMENDATIONS

---

### FOR PRACTITIONERS:

- Use chunked approach for >16 variable Boolean minimization
- Deploy 4+ cores for reasonable turnaround time
- Consider distributed computing for 28+ bit problems
- Verify Verilog output with simulation before synthesis

### FOR RESEARCHERS:

- Investigate cross-chunk optimization opportunities
- Explore alternative chunking strategies (overlap, hierarchy)
- Compare with SAT-based and BDD-based minimizers
- Extend to POS form and multi-output functions

## FUTURE WORK

---

- Benchmark 28-bit and 32-bit functions
- Implement distributed computing across multiple machines
- Optimize chunk processing with GPU acceleration
- Develop interactive visualization for large expressions

## REPRODUCIBILITY

---

Random seed: 42  
Configuration: 24 variables, 4 cores, 256 chunks  
Output files: CSV (chunks + merged), PDF report, Verilog module

Repository: [github.com/Stanislus29/stangorithms](https://github.com/Stanislus29/stangorithms)

All results are reproducible with documented configuration.